

# Lab 4: Cloud Data

## Stat 215A, Fall 2014

Russell Chen, Sreeta Gorripathy, Mo Zhou

November 12, 2014

## 1 Introduction

In this report, we explore cloud detection in the polar regions based on radiances of three images recorded automatically by the MISR sensor aboard the NASA satellite Terra. The classification of cloud and non-cloud is a difficult problem to solve in snow covered regions, due to similarity of cloud and snow cover. In this report, we explore the data using exploration techniques such as k-means, random forests and density plots of the variables. After getting a good idea of the behavior of the image data, we select the important features to create a classification model to identify clouds from non-clouds. For the purpose of classification, we only consider the cloud and non-clouds and discard the not-sure data points. We test out our models against the ?expert labels? for cloud classification to assess the prediction.

The data we have consists of three images from the MISR sensor. For each image, we have the x,y coordinates of pixels, the expert label of the pixels, the radiances from five different angles and three features associated with the pixels, namely NDAI, SD and CORR. The CORR feature is developed by measuring the correlation across the pixels from different angles of image over a patch of the image. The SD feature measures the standard deviation of the radiation from the five angles over a patch of image. NDAI measures the gradient across different patches of the image. The main point that we exploit here is that for non-clouds the same point is captured from different angles and thus the radiation from different angles is similar. We aim to use these information to develop a classification model that can predict for cloud.

## 2 Exploratory Data Analysis

### 2.1 Plots of expert labels for the presence or absence of clouds




Figure 1: The expert labels of the three images. 1 indicates cloud, -1 indicates no cloud, 0 indicates not sure.

There are clustering of cloud and non-cloud for the three images, i.e. the pixels next to some cloud pixel tend to be cloud as well. We also observe that image3 looks different from the other images because it has several small regions of cloud/ non-cloud.

## 2.2 Radiances of Different Angles

The radiances of different angles are similar among the three images. Here we use image1 as an example. The plots below show that the radiances of different angles do not indicate significant differences. The general shapes of the plots are alike and there is no clear difference visually. Density plots of radiances of different angles colored by labels are shown on the next page.




Figure 2: Radiances of different angles of image1

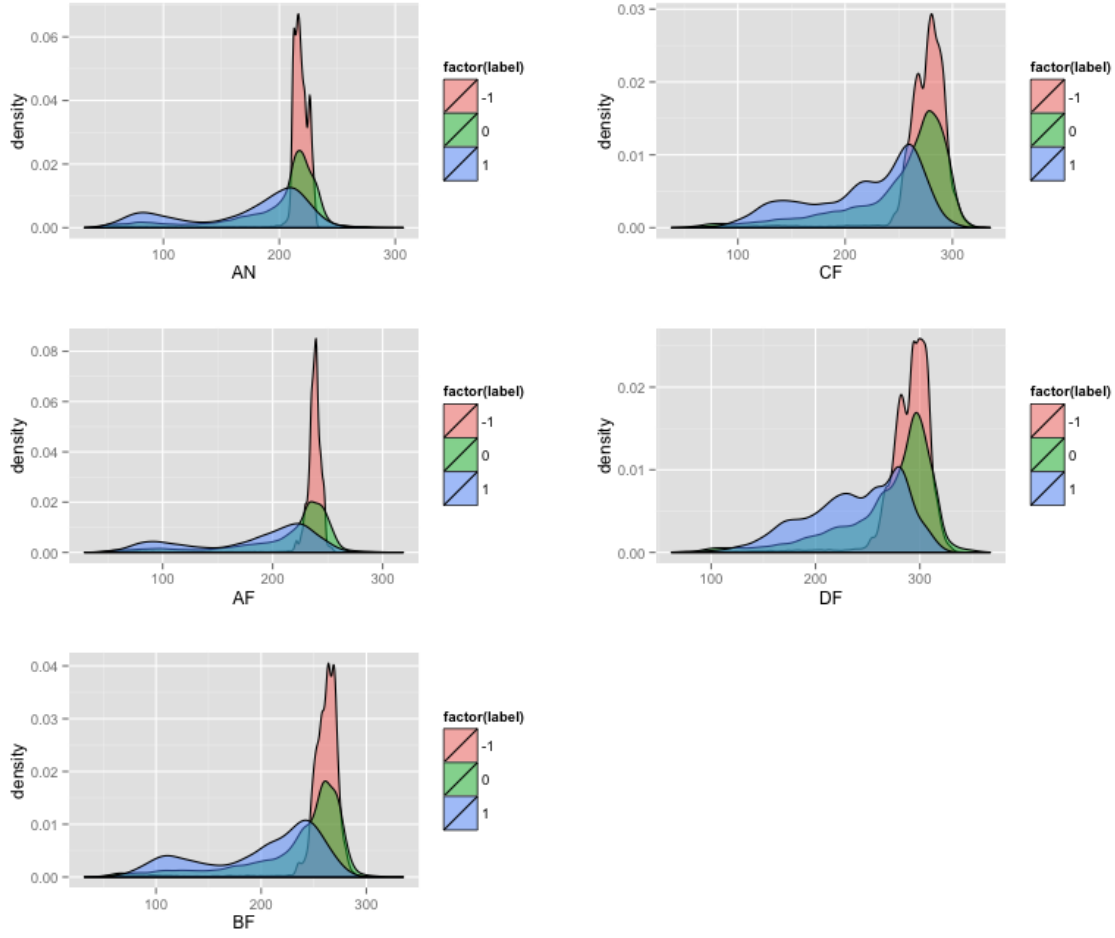


Figure 3: Densities plot of labels with radiances of different angles.

There are no clear intervals that split the labels based on the radiance variables. Thus it is hard to determine if some pixel is cloud just based on the radiance information from different angles. However, we see that the three features provides us better information in classifying the cloud information from non-cloud. The density plots are shown below:




Figure 4: Densities plot of labels with features NDAI, CORR, SD respectively

The density plot of NDAI shows a clear difference in values associated with different labels, which indicates a big difference based on NDAI and suggests that NDAI would be a useful feature to include when deciding

the classification of cloud. The differences based on CORR and SD are not as evident as NDAI, but are still more informative than the radiances, so we might want to consider these two features as well in the classification models.

### 2.3 Unsure labeling treatment

In the images, all pixels are classified by the expert into three groups: cloud, non-cloud and unsure. The pixels with not sure labels are not informative in our classification of cloud, nor is it valuable for us to predict for them. Thus in our analysis and models below, we disregard all the pixels with unsure labeling and are only concerned with the pixels with clear cloud/non-cloud labeling.

## 3 Feature selection

### 3.1 Three best features by density plot

As shown in our exploratory data analysis, the density plot provides us valuable information about the different feature values corresponding to the labels. It shows that NDAI demonstrates the most clear split of the two classes and would be the best feature to consider in the classification models. CORR and SD are important features because some interval of these observations belongs to some particular class. Thus we propose that NDAI, CORR and SD are the best features we wish to include in our classification models.

### 3.2 K-Means clustering based on the three features

To confirm how well the features could predict the classes of our data, we carry out K-Means clustering method on the three images and the three features respectively to obtain the following True Positive Rate (TPR) and False Positive Rate (FPR) results:

		TPR	FPR
Image 1	NDAI	0.9775292	0.1112675
	CORR	0.2122026	0.01879237
	SD	0.2648136	0.03475003
Image 2	NDAI	0.9943717	0.1236416
	CORR	0.4620028	0.009304603
	SD	0.1918963	0.02663122
Image 3	NDAI	0.9416306	0.2492889
	CORR	0.5050838	0.05211543
	SD	0.2305592	0.05288575

### 3.3 Random forest for feature selection

After discarding all the 'unsure' labels, we combined the 3 images and fit a random forest to the data. This allowed us to rank the variables in order of importance for prediction. The `randomForest` package in R calculates variable importance in 2 ways and plots the results, as shown.

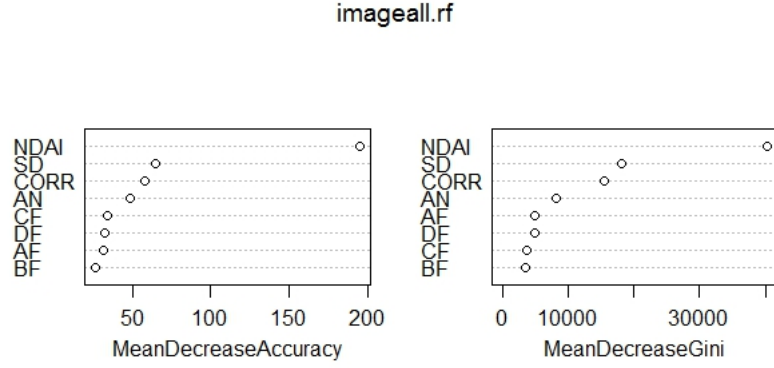


Figure 5: Variable Importance plots generated by Random Forest

On the left, the variables are ranked according to a permutation importance measure. To compute this importance measure for a variable  $p$ , for each tree grown in the random forest, the classification error rate is computed for out-of-bag (OOB) data points in two cases. The first case is where all the variables in the OOB data points are left untouched. The other case is where the values of variable  $p$  are randomly permuted. The classification error rate on the OOB data points for these two cases is compared and the difference between them averaged across all trees to compute a permutation importance measure. A larger difference in the respective classification error rates indicates a variable that has larger predictive power and is therefore more important.

On the right, the variables are ranked according to Gini importance. The Gini index is a measure of dispersion for a list of numbers. Smaller values of the Gini index indicate greater equality and vice versa. In a decision tree, for each split on variable  $p$ , the Gini index for each descendent node is lower than that of the parent node. The larger the Gini index decreases for each split, the better variable  $p$  is at separating the classes and therefore more important for classification.

The two plots shown consistently rank NDAI, SD and CORR as the 3 most important variables for the purposes of prediction. This is also consistent with our earlier reasoning about the importance of NDAI, SD and CORR.

## 4 Classification

### 4.1 Linear Discriminant Analysis (LDA)

LDA models each class (cloud or non-cloud) as a multivariate Normal distribution of the predictor variables and assumes that both classes have the same covariance matrix  $\Sigma$ . These are the class-conditional densities.

$$f_k(x) = \frac{1}{(2\pi)^{p/2} |\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu_k)^T \Sigma^{-1}(x-\mu_k)}$$

Establishing prior probabilities on the classes then allows us to model the posterior probabilities of the classes given the predictor variables in the data by applying Bayes rule. The parameters  $\mu_k$  and  $\Sigma$  of both Gaussian distributions are estimated from training data. In this case, the discriminant functions - and therefore the decision boundaries - are linear in the predictors:

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$$

For the LDA models we fitted, the default class priors  $\pi_k$  were used: empirical proportion of each class in the training set.

## 4.2 Quadratic Discriminant Analysis (QDA)

QDA works in the same fashion as LDA. The class-conditional densities are multivariate Normal distributions of the predictor variables. However, unlike LDA, it does not assume that both classes have the same covariance matrix. The discriminant functions do not simplify; they are quadratic in the predictors:

$$\delta_k(x) = -\frac{1}{2} \log |\Sigma_k| - \frac{1}{2} (x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k) + \log \pi_k$$

Parameters  $\mu_k$  and  $\Sigma_k$  of both Gaussians are also estimated from training data. Again, we used the default class priors due to not having any expert knowledge of cloud cover in the Arctic.

For both LDA and QDA, we did not check the assumptions that the data is Gaussian because we are interested in prediction instead of inference. LDA and QDA models may still do reasonably well at prediction despite the data not being strictly Gaussian. In any case, we will assess all the classification models with the Area under the ROC curve (AUC), as discussed later.

## 4.3 Logistic regression

Since we only have two classes, the logistic regression model consists of only a single linear function:

$$\log \frac{P(\text{label} = \text{no cloud} | X = x)}{1 - P(\text{label} = \text{no cloud} | X = x)} = \beta_0 + \beta_1^T x$$

It is fitted by maximum likelihood in R and the probabilities of each class given the data can be calculated.

## 4.4 Semi-supervised Expectation-Maximization (EM)

An expectation-maximization (EM) algorithm is an iterative method for finding maximum likelihood or maximum a posteriori (MAP) estimates of parameters in statistical models, where the model depends on unobserved latent variables. Semi-supervised EM has proportion of the latent variables known to us, but the other proportion unknown to predict. In our problem, we use the labels of the training set as known latent variables to predict the labels of the testing set. The predictors that are concerned in the Semi-supervised EM model are NDAI, CORR and SD. Other combinations of the variables are also tested, but the performance is not as good.

## 4.5 Support Vector Machine (SVM)

SVMs are supervised learning models in which a hyper-plane is constructed to divide the data into its classes. The training dataset is used to tune its hyper-parameters (gamma and cost) using cross-validation. From experimentation with the sample sizes of the training dataset, we observed that SVM is a very computationally expensive classification method. For our given dataset of 3 images, training on two or one image did not converge even after running for 2 days.

To overcome this problem, we constructed our training dataset by sampling points in the images meant for training. It is interesting to note that the very reason why it is not a good practice to choose points randomly from the images for cross-validation, validates our sampling approach for SVM. In the cloud identification problem, since neighboring pixels are correlated, we gain sufficient information about the training images, by sampling a smaller number of pixels for hyper-parameter tuning.

The hyper-parameters and kernel for the SVM models are chosen from grid-search and tuning on the training set. The SVM model is trained on varying sizes of training data, ranging from 5000 pixels to 15000 pixels.

## 5 Model Assessment

### 5.1 Cross-validation

The images we are dealing with are examples of spatially dependent data. The features of pixels very close to each other are likely to be very similar. The usual method of doing  $k$ -fold cross validation by partitioning the data into  $k$  folds uniformly at random is not ideal since it operates on the assumption that the data points are independent and data points are exchangeable. Thus, when we randomly sample points from all our images in cross-validation, we essentially capture the behavior of larger chunks of data from all the images due to the correlation of pixels. Essentially, this translates to our model 'seeing' all the images and our prediction error is underestimated when we use our classification model on the remaining points of the three images. However, when we apply our prediction on a new image that our training model has not seen before, we will end up getting large prediction errors and poor future prediction.

To overcome this problem, we would like to cross validate in a way that ensures that no data points in folds used for training are close to data points in the folds used for testing. We briefly considered dividing each image into 4 patches by cutting at the midpoint of each of the two coordinate axes and making each patch a single fold in our  $k$ -fold cross validation procedure. However, we observed that training on 11 patches and testing on 1 patch gave extremely unstable AUCs, with AUCs being very low for 2-3 folds and very high for majority of the folds. This might be accounted for the fact that the 12th patch is an 'unseen' image and in some cases, our classification model was not good for the 12th patch. We experimented with the number of patches in the testing set to overcome this problem, but this resulted in an extremely large number of folds leading to expensive computation and the resulting AUC values are not as good as putting 2 images in the training set and the other image in the testing set.

After comparing the model performance for different number of patches, we decided to train on 2 images and then test on the other one. We believe that this is the best simulation of what happens when the classifier is used to predict on a completely new image, since the entire third image is 'unseen' by our classification model.

### 5.2 Akaike Information Criterion (AIC)

The value of the AIC is calculated as:

$$AIC = 2k - 2\log(L)$$

where  $k$  is the number of parameters of the model and  $L$  is the maximized value of the likelihood function for the model. AIC deals with the trade-off between the goodness of fit of the model and the complexity of the model. It is founded on information theory: it offers a relative estimate of the information lost when a given model is used to represent the process that generates the data. We use AIC to assess the logistic regression models with varied number of variables, i.e. with only one variable, three variables, or with four variables. AIC cannot be applied to LDA, QDA because there is no likelihood associated with those methods. By applying AIC to each fold in CV, we summarized results in the following table:

Predictors	Fold	AIC
NDAI, CORR, SD	1	72739.29
	2	81264.35
	3	72984.58
	joint	116338
NDAI	1	79758.82
	2	85215.58
	3	92612.87
	joint	130234.2
NDAI, CORR, SD, AN	1	68693.97
	2	81164.8
	3	71237.18
	joint	114806.1

Since a lower AIC indicates a better model, we can see that the third set of 4 variables with NDAI, CORR, SD and AN is the best logistic model among the three. Notice also that the AIC values for 3 variables with NDAI, CORR and SD are similar to the model including AN.

### 5.3 Receiver Operating Characteristic (ROC) curves

We use ROC to assess the stability of the models with disturbance in threshold. Each model returns a vector of the probability of a particular pixel being in class 1. By deciding classification using different thresholds, we obtain the ROC curve that summarizes the result. We also compute the area under the ROC curve (AUC) to assess the performance of the models. The AUC values are shown in the tables below.

Predictors	Fold	AUC for logit	Predictors	Fold	AUC for LDA
NDAI, CORR, SD	1	0.8970	NDAI, CORR, SD	1	0.8953
	2	0.9554		2	0.9543
	3	0.9356		3	0.9512
	average	<b>0.9294</b>		average	<b>0.9336</b>
	joint	<b>0.9396</b>		joint	<b>0.9397</b>
NDAI	1	0.8813	NDAI	1	0.8813
	2	0.9422		2	0.9422
	3	0.9571		3	0.9571
	average	<b>0.9269</b>		average	<b>0.9269</b>
	joint	<b>0.9305</b>		joint	<b>0.9295</b>
NDAI, CORR, SD, AN	1	0.8988	NDAI, CORR, SD, AN	1	0.8988
	2	0.9521		2	0.9485
	3	0.9319		3	0.9422
	average	<b>0.9276</b>		average	<b>0.9298</b>
	joint	<b>0.9330</b>		joint	<b>0.9307</b>

Predictors	Fold	AUC for QDA
NDAI, CORR, SD	1	0.8878
	2	0.9707
	3	0.9594
	average	<b>0.9393</b>
	joint	<b>0.9473</b>
NDAI	1	0.8818
	2	0.9422
	3	0.9596
	average	<b>0.9279</b>
	joint	<b>0.9302</b>
NDAI, CORR, SD, AN	1	0.8765
	2	0.9607
	3	0.9532
	average	<b>0.9301</b>
	joint	<b>0.9403</b>

Predictors	Fold	AUC for E-M
NDAI, CORR, SD	1	0.8886
	2	0.9600
	3	0.9541
	average	<b>0.9342</b>
	joint	<b>0.9260</b>

As shown in the above tables, QDA has the highest joint AUC and average AUC values. However, QDA does not predict as good for image3 as logistic regression does. LDA and QDA have similar AUCs for each fold and QDA is predicting slightly better. EM gives lower AUC than QDA and the prediction for image3 is not as good as logistic regression. The overall performance of the 4 models are similar with logistic regression being the most stable and QDA with the best performance.

The ROC curves for the best QDA model (3 predictors - NDAI, SD and CORR) are plotted below. The top left plot is for fold 1, the top right plot for fold 2, the bottom left plot for fold 3 and the bottom right plot shows the pooled ROC curve. The ROC curves for the rest of the models are very similar to the 4 shown



here. The 4 plots shown here are close to ideal since they cover a large area of the plot and approach the top left point closely. This top left point is the “perfect” point where the TPR is 1 and the FPR is 0. The shape of the plots are very similar across folds, which indicates that the prediction behaves in a stable manner. This prediction stability gives us confidence in our classifier for future prediction of new images.




Figure 6: ROC curves for QDA with NDAI, SD and CORR

#### 5.4 Model Assessment for SVM classifier

The measures to validate the classification for different training sizes are plotted in the figure below. It can be seen that across different measures such as accuracy ( $((TP+TN)/(TP + TN + FP + FN))$ ), TPR ( $(TP/(TP + FN))$ ), FPR ( $(FP/(FP + TN))$ ) and FDR ( $(FP/(FP + TP))$ ), SVM performs poorly as a classifier for the cloud data.

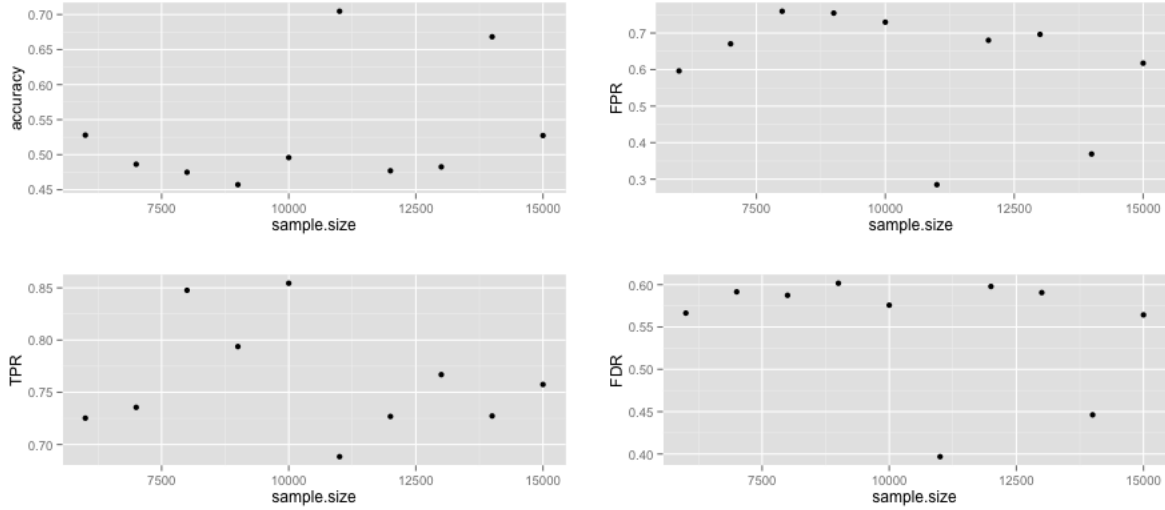


Figure 7: Model assessment for SVM

It is surprising that these measures do not show a consistent trend as the sample size increases, which may be accounted to the fact that the strength of this classifier is largely dependent on how representative the training set is. Since the sampling of the training set is random, we can never be sure that the training data has a representative sample of for both cloud and non-cloud. Since increasing the sample size doesn't directly address this problem, we do not see consistent trends of improvement in performance due to increase in training sample size. Thus we do not consider the SVM classifier in its current formulation as an appropriate model for classifying cloud and non-cloud.

## 6 Best Classification Model

A closer look at the AUC from the logistic model, LDA and QDA makes us think that there might be a better model that combines these two models and gives us more stable prediction. Notice that QDA seems to have the best performance by having the greatest average AUC as well as joint AUC. However, QDA gives the worst AUC when predicting for image3. In addition, logistic regression provides the best AUC when predicting for image 3 but does not perform well for image 3. Thus we propose that a more stable model can be achieved by giving weights to the probability vectors predicted from the logistic model and QDA. We do not include LDA because the performance of LDA is not as good when predicting all three images.

A naive model to consider is to weigh the probability vectors the same, i.e. 0.5 on the logistic model prediction probability and 0.5 on the QDA logistic model prediction. We used this model and assess the fit for the three folds. The AUC for the three folds are 0.8949305, 0.9639261 and 0.9573461 respectively. The average AUC is slightly smaller than the average AUC of the best QDA model but the variance of AUC is smaller in this model, which demonstrate an improvement in stability. Changing the weights on the two models does not change the result too much since the AUC for both methods are similar to begin with. Hence we will just work with the ensemble model with equal weights on the probability vectors from the logistic model and QDA.

We would use this model as our best model for predictions on new images by training on all three images we have now.

## 7 Misclassification Error Analysis

In this section, we explore the misclassification error from the best classification model. The following plots show the misclassification locations, expert labels and the original image of image1. Note that the same pattern is observed in the misclassification of the other two images as well, so we show the plots for image3 as an example since image3 has the biggest misclassification error.




Figure 8: The left plot shows the pixels that are misclassified. The middle plot shows the original expert labels. The right plot shows the actual image from AN radiance.

As we can see, the misclassification locations tend to cluster in particular regions of the image. Misclassification happens when the region is small and distant from the other regions with the same label. Furthermore, misclassification locations are scattered at the pixels with bigger variations in the region. Those variations may be the landscape of the location.

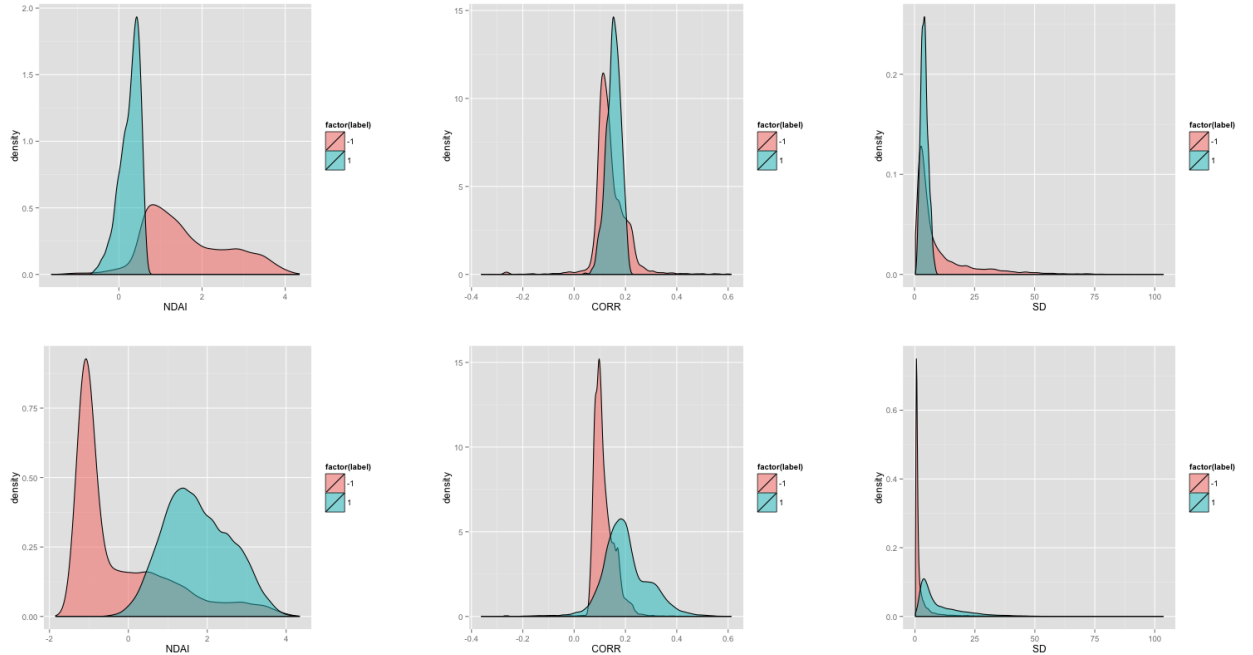


Figure 9: The first row is the density plots of NDAI, CORR and SD of the misclassification errors. The second row is the original density plots of NDAI, CORR and SD of the same image.

By the density plots of the misclassification errors, we can see that misclassification mainly happens in the

intersection of cloud/no cloud in the three featured values. This is reasonable because when our model tries to find the probability of classification, the intersection of values in the features are the hardest to separate. Hence misclassification mainly occurs in those regions where the featured values are similar but the labels are different.

## 8 Future Prediction

Future prediction is a very important aspect for choosing the best classifier model and is often ignored. The best models are chosen based on the available data (cross-validating on training and test) using measures such as AUC, AIC and BIC. However, it is important to remember that these measures indicate the accuracy of the model in predicting for the type of data that is available before-hand. Future prediction may be bad, if the image is very different from the data that the classification was trained on.

To account for such situations, we focused on the stability of AUC across the different folds during our model selection. The average AUC of our best classification model was slightly less than the average AUC of its component QDA model. The better stability of the AUC over different folds in the best classification model shows that the ensemble model works better on new or different images than the QDA model. Since future prediction is important for our classifier, we choose the ensemble as our best model. Thus, our classification model is poised to work well on future data without expert labels.

In addition, our approach of choosing the cross-validation folds ensures good future prediction. Since we trained our classification model on two images and tested on the third, we have ensured that our model is being tested on 'unseen' image and data points. Our best model has been chosen from the performance of the models on the unseen images in 3 folds and simulates the future prediction situations in which we are given completely new images unseen by the classifier.

## References

[1] <https://github.com/mozhou/cloud-data-lab>